Jack Sturtevant
jack@sturt.io

Hello AssemblyAI Team,

I'm applying for the Senior Software Engineer, Inference role. Taking a model that works in a notebook and making it fast and dependable for users is the part of AI work I enjoy most.

At Coordinator I fine-tuned multi-sport LLMs on H100s and served them quantized with vLLM. I also benchmarked self-hosted Whisper and TTS engines before shipping voice messaging in our web app, and built a VLM video pipeline with a scoring harness so we could tune thresholds against real footage. When our quick-edit path was slow, I timed a single-shot model call against the agent tiers on real requests. It was 4–6x faster, so I shipped it as its own tier. Earlier, as founder and CTO of Ralli, I built NestJS and Postgres services with live WebSocket updates and Datadog tracing.

I haven't run inference at the scale of a million API users, and I'd be learning that from your team. I'd be glad to talk about how I could help get research into customers' hands.

Jack Sturtevant