Jack Sturtevant
jack@sturt.io

Hello Hugging Face Team,

I'm applying for the Open-Source Machine Learning Engineer role. Almost everything I built this year ran on open models, and much of it ran because of the Hugging Face stack, so working on that stack is the most direct way I can think of to pay it back.

At Coordinator I fine-tuned LLMs on H100s and served them quantized with vLLM, ran Qwen3-VL and Cosmos-Reason2 inference to detect football formations from game film, and benchmarked self-hosted Whisper and TTS engines before shipping voice messaging. Each of those started with a model card, a pinned revision, and a lot of reading other people's code to find out why something behaved the way it did. I want to be on the side of that work that makes the next person's afternoon shorter.

To be direct about the gap: my public open-source work so far is compiler work on Dafny from an AWS internship, not merged pull requests to ML libraries, and my PyTorch use has been as a practitioner rather than a library author. I'd be glad to talk about whether that fits, and I'd start contributing either way.

Jack Sturtevant