Jack Sturtevant
jack@sturt.io

Hello Pictor Labs Team,

I'm applying for the Sr. Machine Learning Engineer role because I want my GPU inference work to help patients. I started college as an EMT studying biomedical engineering, and my first job was in a Harvard medical lab working on early lab-on-a-chip experiments. Coding pulled me away, but I've kept taking biology and chemistry classes on my own time and money.

At Coordinator I fine-tuned LLMs on H100s, wrote the quantization to serve them with vLLM, and built computer vision pipelines over football film. I also moved our GPU voice services to ARM64 NVIDIA hardware, which is close to the edge devices you deploy on. When a path was slow, I timed it on real requests, and one change I shipped that way was 4-6x faster. Before that, I wrote C++ for robotics and machine learning down to the CUDA level, though that was about five years ago.

I don't have 7+ years in production inference or 5+ years of PyTorch. I also haven't used TensorRT or Triton Inference Server in production, or worked with gigapixel pathology images. I'd be learning those from your team.

I'd be glad to talk about how I could help get whole slide images through in under 2 minutes.

Jack Sturtevant