Jack Sturtevant
jack@sturt.io

Hello Twilio Customer Memory Team,

I'm applying for the Machine Learning Engineer 2 role on Customer Memory. Memory that an agent can actually rely on is the problem I spent the most time on this year.

At Coordinator I built the agent runtime in Python, including group-scoped memory pruning and voice-agent memory, and a pipeline that turned reviewed production turns into training data for fine-tuned models served on vLLM. I measured instead of guessing: I tuned vision-model thresholds against a scoring harness, and timed single-shot against agent-loop tiers on real requests (single-shot was 4-6x faster, so it became a fourth tier). I wrote the Go agent tools with exact contracts, and I work agentic-first with Claude Code and Cursor.

To be direct about gaps: I haven't used Kafka or another event-streaming system in production, and my retrieval work has been at startup scale rather than Twilio's.

I'd be glad to talk about the role.

Jack Sturtevant